

01_churn_prediction

October 23, 2025

1 Task 1: Customer Churn Prediction for a Telecom Company

Task Overview Objective: Build a machine learning model to predict customer churn using historical data.

Deliverables:

- Exploratory Data Analysis (EDA)
- Feature engineering
- Train/test split and model selection (Logistic Regression, XGBoost, etc.)
- Performance metrics (confusion matrix, AUC-ROC)
- Final report with visualizations

[205]: # Mock Data (Python):

```
import pandas as pd
import numpy as np

np.random.seed(42)

n = 10000

data = pd.DataFrame({
    'CustomerID': np.arange(n),
    'Gender': np.random.choice(['Male', 'Female'], size=n),
    'SeniorCitizen': np.random.choice([0, 1], size=n),
    'Tenure': np.random.randint(1, 72, size=n),
    'MonthlyCharges': np.round(np.random.uniform(20, 120, size=n), 2),
    'TotalCharges': lambda df: df['Tenure'] * df['MonthlyCharges'],
    'Contract': np.random.choice(['Month-to-month', 'One year', 'Two year'], size=n),
    'PaymentMethod': np.random.choice(['Electronic check', 'Mailed check', 'Bank transfer', 'Credit card'], size=n),
    'Churn': np.random.choice([0, 1], size=n, p=[0.73, 0.27])
})

data['TotalCharges'] = (data['Tenure'] * data['MonthlyCharges']).round(2)
```

[206]: data.head()

```
[206]:   CustomerID  Gender  SeniorCitizen  Tenure  MonthlyCharges  TotalCharges \
0          0    Male           0      55       111.88     6153.40
1          1  Female          1      36        58.70     2113.20
2          2    Male           0      37       118.86     4397.82
3          3    Male           1      14        96.14     1345.96
4          4    Male           1      27       28.05      757.35

          Contract  PaymentMethod  Churn
0  Two year    Mailed check     0
1  Two year  Electronic check     0
2  One year  Electronic check     0
3 Month-to-month    Mailed check     1
4  Two year    Mailed check     0
```

```
[207]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 9 columns):
 #   Column            Non-Null Count  Dtype  
--- 
 0   CustomerID        10000 non-null   int64  
 1   Gender             10000 non-null   object  
 2   SeniorCitizen      10000 non-null   int64  
 3   Tenure             10000 non-null   int64  
 4   MonthlyCharges    10000 non-null   float64 
 5   TotalCharges       10000 non-null   float64 
 6   Contract           10000 non-null   object  
 7   PaymentMethod      10000 non-null   object  
 8   Churn              10000 non-null   int64  
dtypes: float64(2), int64(4), object(3)
memory usage: 703.2+ KB
```

```
[216]: data['Churn'].value_counts()
```

```
[216]: Churn
0    7330
1    2670
Name: count, dtype: int64
```

```
[209]: X = data[['Gender', 'SeniorCitizen', 'Tenure', 'MonthlyCharges', 'TotalCharges', 'Contract', 'PaymentMethod']]
y = data['Churn']
```

```
[ ]: from sklearn.preprocessing import LabelEncoder, StandardScaler
le = LabelEncoder()
```

```

X.loc[:, 'Gender'] = le.fit_transform(X['Gender'])
X.loc[:, 'Contract'] = le.fit_transform(X['Contract'])
X.loc[:, 'PaymentMethod'] = le.fit_transform(X['PaymentMethod'])

# preprocess the data
scalar = StandardScaler()
X.loc[:, ['Tenure', 'MonthlyCharges', 'TotalCharges']] = scalar.
    fit_transform(X[['Tenure', 'MonthlyCharges', 'TotalCharges']])
X.loc[:, ['Contract', 'PaymentMethod']] = scalar.fit_transform(X[['Contract', u
    'PaymentMethod']])
```

[211]: X.sample(n=5)

	Gender	SeniorCitizen	Tenure	MonthlyCharges	TotalCharges	Contract	\
1780	0		1 -0.339256	-1.705627	-1.026792	-0.013227	
4401	1		0 0.099753	-0.153487	-0.017785	-0.013227	
5129	0		0 -0.436814	0.388431	-0.178860	-0.013227	
3968	0		1 -0.534371	-0.649094	-0.665083	-1.249411	
9678	0		1 -1.119716	-0.761763	-1.017518	-1.249411	
			PaymentMethod				
1780			-1.340355				
4401			1.344651				
5129			0.449649				
3968			0.449649				
9678			0.449649				

[212]: from sklearn.model_selection import train_test_split

```

# train test split the data
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,u
    random_state=42, stratify=y)
```

[213]: from sklearn.neighbors import KNeighborsClassifier

```

model = KNeighborsClassifier(n_neighbors=8)
model.fit(X_train, y_train)

# predict the model
y_pred = model.predict(X_test)
```

[214]: import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.metrics import roc_curve, roc_auc_score

Predict probabilities
y_probs = model.predict_proba(X_test)

```

print("Probabilities of being True Negative and True Positive: ")
print(y_probs[:5])

y_pred_probs = y_probs[:, 1] # Picked 2nd column which contains
    ↪Predicted-Yes(True Positive) values. and 1st column contains
    ↪Predicted-No(True Negative) values
print("Probabilities of being True Positive: ", y_pred_probs[:5])

# Generate ROC curve values: fpr, tpr, thresholds
fpr, tpr, thresholds = roc_curve(y_test, y_pred_probs)
plt.plot([0, 1], [0, 1], 'k--')

# Plot tpr against fpr
plt.plot(fpr, tpr)
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve for Diabetes Prediction')
plt.show()

#####
print(roc_auc_score(y_test, y_pred_probs)) # Calculate roc_auc_score

```

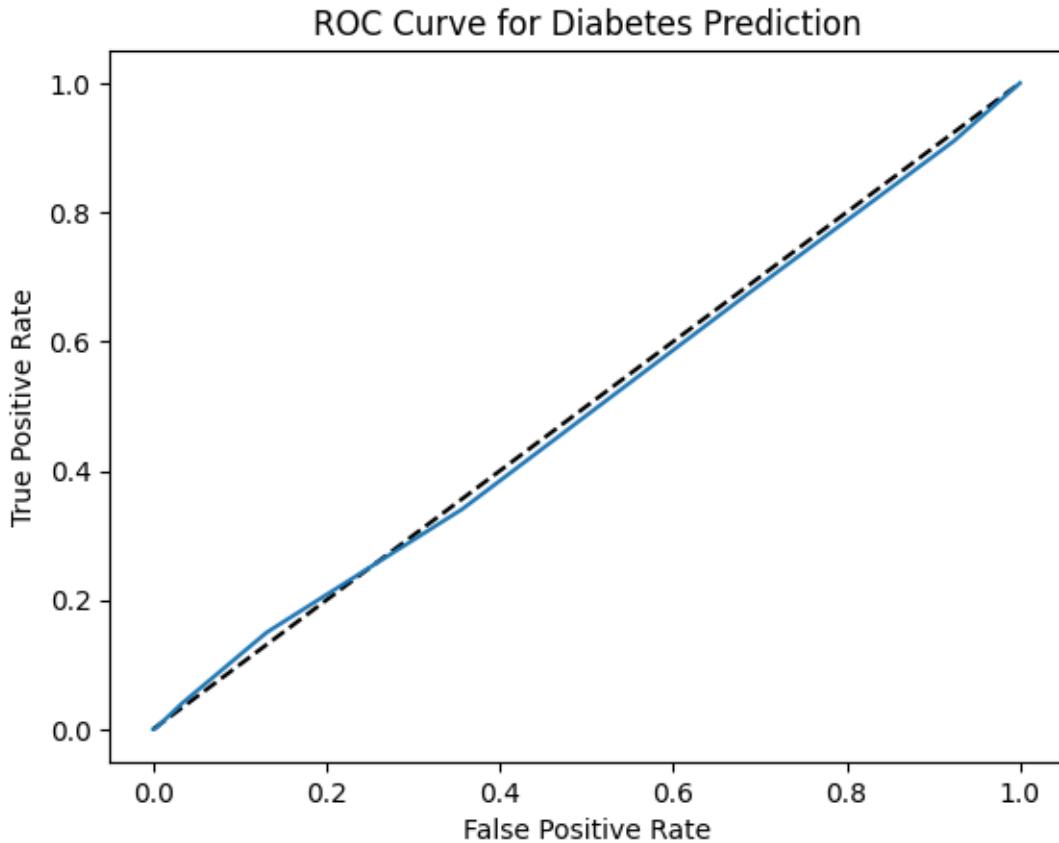
Probabilities of being True Negative and True Positive:

```

[[0.875 0.125]
 [0.875 0.125]
 [0.75  0.25 ]
 [0.625 0.375]
 [0.75  0.25 ]]

```

Probabilities of being True Positive: [0.125 0.125 0.25 0.375 0.25]



0.4934373898247927

```
[215]: from sklearn.metrics import confusion_matrix, classification_report

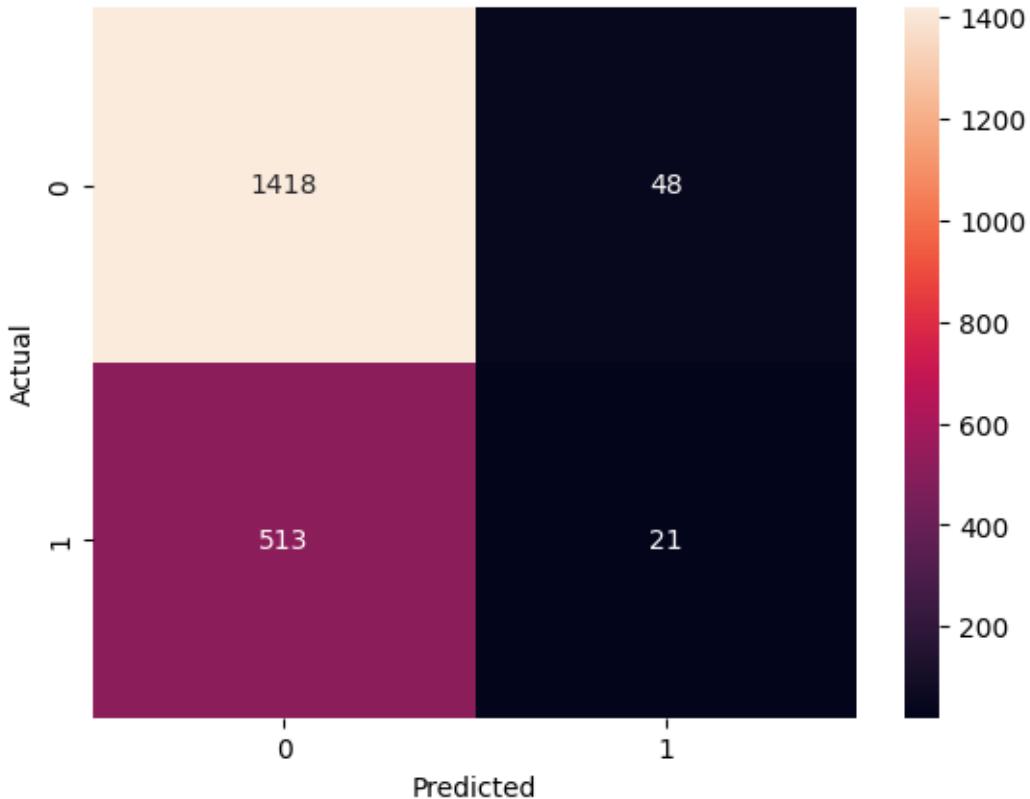
# evaluate the model
print(confusion_matrix(y_test, y_pred))
print(classification_report(y_test, y_pred))

# plot confusion matrix
sns.heatmap(confusion_matrix(y_test, y_pred), annot=True, fmt='d')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.show()
```

```
[[1418  48]
 [ 513  21]]
```

	precision	recall	f1-score	support
0	0.73	0.97	0.83	1466
1	0.30	0.04	0.07	534

accuracy			0.72	2000
macro avg	0.52	0.50	0.45	2000
weighted avg	0.62	0.72	0.63	2000



1.1 REPORT:

The model is evaluated on a dataset of 2000 samples. The results are as follows:

- Our model is predicting 72% accurately.
- The model performs well on **Class 0**, achieving high recall (0.97), meaning it correctly identifies most of the **Class 0** samples.
- For **Class 1**, performance is very poor, with recall only 0.04. This shows the model is failing to detect **Class 1** cases.

The model is biased toward **Class 0** and struggles with detecting **Class 1**