

# 03\_sales\_forecasting

October 23, 2025

## 1 Task 3: Sales Forecasting Using Time Series Analysis

### Task Overview

Objective: Forecast the next 12 months of product sales using time series models.

### Deliverables:

- Time series decomposition (trend, seasonality)
- ARIMA, Prophet or LSTM modeling
- Evaluation metrics (MAE, RMSE)
- Forecast visualizations

```
[41]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from statsmodels.tsa.seasonal import seasonal_decompose
```

```
[42]: # Mock Data (Python):
date_range = pd.date_range(start='2020-01-01', periods=36, freq='M')
sales = np.random.poisson(lam=200, size=36) + np.linspace(0, 100, 36)
df = pd.DataFrame({'Date': date_range, 'Sales': sales.round(2)})
```

/tmp/ipykernel\_478162/2065811994.py:2: FutureWarning:

'M' is deprecated and will be removed in a future version, please use 'ME' instead.

```
[43]: df.head(10)
```

```
[43]:      Date    Sales
0 2020-01-31  182.00
1 2020-02-29  223.86
2 2020-03-31  200.71
3 2020-04-30  213.57
4 2020-05-31  220.43
5 2020-06-30  208.29
```

```
6 2020-07-31 214.14
7 2020-08-31 208.00
8 2020-09-30 219.86
9 2020-10-31 220.71
```

```
[44]: df.info()
```


```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 36 entries, 0 to 35
Data columns (total 2 columns):
 #   Column    Non-Null Count  Dtype  
--- 
 0   Date      36 non-null    datetime64[ns]
 1   Sales     36 non-null    float64 
dtypes: datetime64[ns](1), float64(1)
memory usage: 704.0 bytes
```

```
[45]: df.describe()
```

```
[45]:
```

	Date	Sales
count	36	36.000000
mean	2021-07-15 21:20:00	248.777778
min	2020-01-31 00:00:00	182.000000
25%	2020-10-23 06:00:00	220.357500
50%	2021-07-15 12:00:00	247.640000
75%	2022-04-07 12:00:00	274.215000
max	2022-12-31 00:00:00	314.570000
std		33.927472

```
[46]: sns.lineplot(data=df, x="Date", y="Sales")
plt.xticks(rotation=45)
plt.show()
```




```
[47]: df['year'] = df['Date'].dt.year  
df['month'] = df['Date'].dt.month
```

```
[48]: print(df['month'].unique())  
print(df.year.unique())
```

```
[ 1  2  3  4  5  6  7  8  9 10 11 12]  
[2020 2021 2022]
```


```
[49]: # draw a line plot for each month  
sns.lineplot(data=df, x="month", y="Sales", hue="year", palette='deep')  
plt.xlim(1, 12)
```


```
[49]: (1.0, 12.0)
```



```
[50]: result = seasonal_decompose(df['Sales'], model='additive', period=12) # period
      ↵for 1 year
      result.plot()
```

```
[50]:
```





```
[51]: # rename the columns
df = df.rename(columns={'Date':'ds', 'Sales': 'y'}) # prophet model take date as ds and target variable as y
df.head()
```

```
[51]:      ds      y  year  month
0 2020-01-31  182.00  2020      1
1 2020-02-29  223.86  2020      2
2 2020-03-31  200.71  2020      3
3 2020-04-30  213.57  2020      4
4 2020-05-31  220.43  2020      5
```

```
[ ]: from prophet import Prophet
from prophet.plot import plot_plotly

model = Prophet() # call the model
model.fit(df)

# make future dataframe and forecast
forecast = model.make_future_dataframe(periods=12, freq='M') # forecast for 12 months
```

```

forecast = model.predict(forecast)

INFO    prophet:forecaster.py:parse_seasonality_args()- Disabling weekly
seasonality. Run prophet with weekly_seasonality=True to override this.
INFO    prophet:forecaster.py:parse_seasonality_args()- Disabling daily
seasonality. Run prophet with daily_seasonality=True to override this.
DEBUG   cmdstanpy:filesystem.py:_temp_single_json()- input tempfile:
/tmp/tmp56bic091/0dc0yy4r.json
DEBUG   cmdstanpy:filesystem.py:_temp_single_json()- input tempfile:
/tmp/tmp56bic091/6w4lfv28.json
DEBUG   cmdstanpy:model.py:_run_cmdstan()- idx 0
DEBUG   cmdstanpy:model.py:_run_cmdstan()- running CmdStan, num_threads: None
DEBUG   cmdstanpy:model.py:_run_cmdstan()- CmdStan args: ['/home/tk-
lpt-648/miniconda3/envs/data_analysis/lib/python3.10/site-
packages/prophet/stan_model/prophet_model.bin', 'random', 'seed=50775', 'data',
'file=/tmp/tmp56bic091/0dc0yy4r.json', 'init=/tmp/tmp56bic091/6w4lfv28.json',
'output',
'file=/tmp/tmp56bic091/prophet_modeli82mjpc1/prophet_model-20251001144505.csv',
'method=optimize', 'algorithm=newton', 'iter=10000']
14:45:05 - cmdstanpy - INFO - Chain [1] start processing
INFO    cmdstanpy:model.py:_run_cmdstan()- Chain [1] start processing
14:45:05 - cmdstanpy - INFO - Chain [1] done processing
INFO    cmdstanpy:model.py:_run_cmdstan()- Chain [1] done processing
/home/tk-lpt-648/miniconda3/envs/data_analysis/lib/python3.10/site-
packages/prophet/forecaster.py:1872: FutureWarning:

'M' is deprecated and will be removed in a future version, please use 'ME'
instead.

```

[66]: forecast.tail()

```

[66]:      ds      trend  yhat_lower  yhat_upper  trend_lower  trend_upper  \
43 2023-08-31  331.879207  350.776320  370.404424  331.663277  332.106804
44 2023-09-30  334.956461  272.094432  291.079138  334.697064  335.229913
45 2023-10-31  338.136291  324.271413  342.306453  337.822754  338.452744
46 2023-11-30  341.213545  329.468711  349.015464  340.859587  341.579157
47 2023-12-31  344.393375  336.810030  356.803024  343.992231  344.797360

      additive_terms  additive_terms_lower  additive_terms_upper  yearly  \
43          28.792010                  28.792010          28.792010  28.792010
44         -53.785283                 -53.785283         -53.785283 -53.785283
45          -4.930613                  -4.930613          -4.930613 -4.930613
46         -1.755094                  -1.755094          -1.755094 -1.755094
47          2.892748                  2.892748          2.892748  2.892748

yearly_lower  yearly_upper  multiplicative_terms  \

```

```

43    28.792010    28.792010          0.0
44   -53.785283   -53.785283          0.0
45   -4.930613   -4.930613          0.0
46   -1.755094   -1.755094          0.0
47    2.892748    2.892748          0.0

      multiplicative_terms_lower  multiplicative_terms_upper      yhat
43                  0.0           0.0  360.671217
44                  0.0           0.0  281.171179
45                  0.0           0.0  333.205678
46                  0.0           0.0  339.458451
47                  0.0           0.0  347.286123

```

[67]: df.shape, forecast.shape # 36 + 12 = 48

[67]: ((36, 4), (48, 16))

```

[68]: # calculate mean squared error
from sklearn.metrics import mean_squared_error, mean_absolute_error
y_true = df['y']
y_pred = forecast['yhat'][:len(y_true)]

print(f'Mean Squared Error: {mean_squared_error(y_true, y_pred):.2f}')
print("MAE: ", mean_absolute_error(y_true, y_pred))
print("RMSE: ", np.sqrt(mean_squared_error(y_true, y_pred)))

```

Mean Squared Error: 54.86

MAE: 6.100577729932693

RMSE: 7.406436716511218

[69]: # plot the forecast  
plotly(model, forecast)