

04_fraud_detection

October 23, 2025

1 Task 4: Fraud Detection in Financial Transactions

1.1 Task Overview

Objective: Identify fraudulent transactions from synthetic financial data.

Deliverables:

- Data preprocessing and visualization
- Class balancing (SMOTE or undersampling)
- Random Forest / Gradient Boosting model
- Confusion matrix, Precision/Recall analysis

```
[ ]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from imblearn.over_sampling import SMOTE
from imblearn.under_sampling import RandomUnderSampler
from imblearn.pipeline import Pipeline
```

```
[ ]: n = 50000
data = pd.DataFrame({
    'TransactionID': np.arange(n),
    'Amount': np.round(np.random.exponential(scale=100, size=n), 2),
    'Time': np.random.randint(0, 86400, size=n), # seconds in a day
    'Location': np.random.choice(['NY', 'CA', 'TX', 'FL'], size=n),
    'Device': np.random.choice(['Mobile', 'Web', 'ATM'], size=n),
    'IsFraud': np.random.choice([0, 1], size=n, p=[0.98, 0.02])
})
```

```
[103]: data.sample(5)
```

```
[103]:   TransactionID  Amount      Time Location  Device  IsFraud
      5548          5548  229.46  19688       FL  Mobile      0
     37543          37543   12.16  24901       TX    Web      0
    22735          22735  143.40  16649       TX  Mobile      0
     3494          3494   59.95  61208       NY    ATM      0
```

```
42802      42802  294.04  83431      TX      Web      0
```

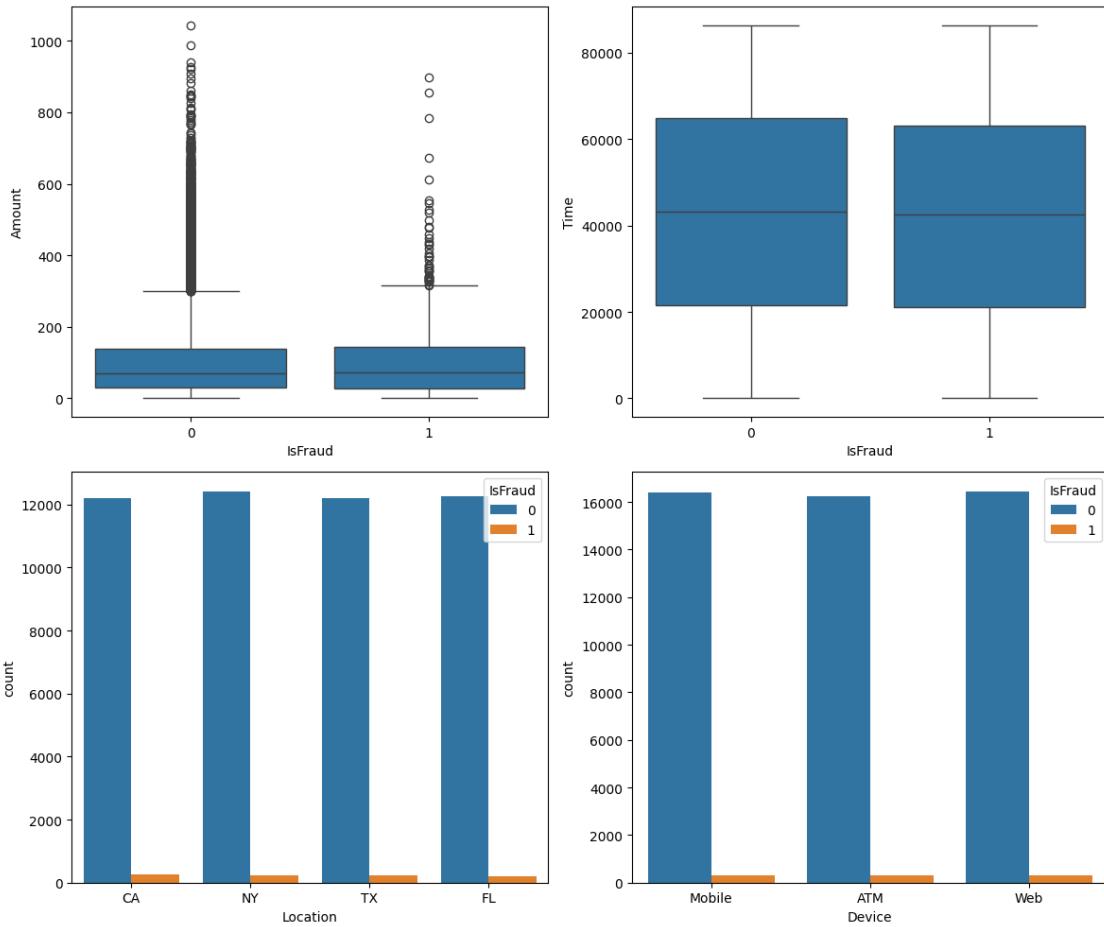
```
[104]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 50000 entries, 0 to 49999
Data columns (total 6 columns):
 #   Column            Non-Null Count  Dtype  
--- 
 0   TransactionID    50000 non-null   int64  
 1   Amount             50000 non-null   float64 
 2   Time               50000 non-null   int64  
 3   Location           50000 non-null   object  
 4   Device              50000 non-null   object  
 5   IsFraud            50000 non-null   int64  
dtypes: float64(1), int64(3), object(2)
memory usage: 2.3+ MB
```

```
[105]: data.describe()
```

```
TransactionID      Amount        Time      IsFraud
count      50000.000000  50000.000000  50000.000000  50000.000000
mean       24999.500000   100.478261  43233.109520   0.018500
std        14433.901067   100.564985  24910.786162   0.134752
min         0.000000    0.000000     0.000000    0.000000
25%        12499.750000   29.305000   21647.750000   0.000000
50%        24999.500000   70.280000   43265.500000   0.000000
75%        37499.250000  138.110000  64926.250000   0.000000
max        49999.000000  1043.960000  86399.000000   1.000000
```

```
[106]: fig, axs = plt.subplots(2, 2, figsize=(12, 10))
sns.boxplot(x='IsFraud', y='Amount', data=data, ax=axs[0, 0])
sns.boxplot(x='IsFraud', y='Time', data=data, ax=axs[0, 1])
sns.countplot(x='Location', hue='IsFraud', data=data, ax=axs[1, 0])
sns.countplot(x='Device', hue='IsFraud', data=data, ax=axs[1, 1])
plt.tight_layout()
plt.show()
```



```
[107]: data.columns
```

```
[107]: Index(['TransactionID', 'Amount', 'Time', 'Location', 'Device', 'IsFraud'],
dtype='object')
```

```
[108]: X = data[['Amount', 'Time', 'Location', 'Device']]
y = data['IsFraud']
```

```
[109]: print("Original class distribution:")
print(data['IsFraud'].value_counts())
```

```
Original class distribution:
IsFraud
0    49075
1     925
Name: count, dtype: int64
```

```
[112]: # Preprocess data
from sklearn.preprocessing import LabelEncoder, MinMaxScaler

le = LabelEncoder()
X.loc[:, 'Location'] = le.fit_transform(X['Location'])
X.loc[:, 'Device'] = le.fit_transform(X['Device'])

# preprocess the data
scalar = MinMaxScaler()
X.loc[:, ['Amount', 'Time']] = scalar.fit_transform(X[['Amount', 'Time']])
```

/tmp/ipykernel_514198/2501154083.py:10: FutureWarning: Setting an item of incompatible dtype is deprecated and will raise in a future error of pandas.
Value '[0.88507969 0.94760356 0.444982 ... 0.7705992 0.69556361 0.48238984]' has dtype incompatible with int64, please explicitly cast to a compatible dtype first.

```
X.loc[:, ['Amount', 'Time']] = scalar.fit_transform(X[['Amount', 'Time']])
```

```
[113]: X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, stratify=y, random_state=42
)
```

```
[114]: # --- Balancing with SMOTE + Undersampling ---
over = SMOTE(sampling_strategy=0.1, random_state=42)
under = RandomUnderSampler(sampling_strategy=0.5, random_state=42)

balancing_pipeline = Pipeline(steps=[
    ('smote', over),
    ('under', under)
])
```

```
[115]: # Apply balancing
X_train_res, y_train_res = balancing_pipeline.fit_resample(X_train, y_train)
```

```
[116]: len(X_train_res), y_train_res.value_counts()
```

```
[116]: (11778,
       IsFraud
       0    7852
       1    3926
Name: count, dtype: int64)
```

```
[ ]: from sklearn.ensemble import RandomForestClassifier

model = RandomForestClassifier(random_state=42,
                               n_estimators=100,
                               # min_samples_leaf=0.05
```

```
)  
model.fit(X_train_res, y_train_res)
```

[]: RandomForestClassifier(random_state=42)

[118]: y_pred = model.predict(X_test)

```
[119]: from sklearn.metrics import accuracy_score, confusion_matrix,  
       classification_report  
  
#evaluate the model  
print('accuracy score: ', accuracy_score(y_test, y_pred))  
print('confusion matrix:\n', confusion_matrix(y_test, y_pred))  
print('classification report:\n', classification_report(y_test, y_pred))
```

```
accuracy score: 0.8498  
confusion matrix:  
[[8477 1338]  
 [ 164   21]]  
classification report:  
          precision    recall  f1-score   support  
  
          0       0.98      0.86      0.92      9815  
          1       0.02      0.11      0.03      185  
  
     accuracy                           0.85      10000  
    macro avg       0.50      0.49      0.47      10000  
weighted avg       0.96      0.85      0.90      10000
```

1.2 REPORT:

The model is evaluated on a dataset of 10000 samples. The results are as follows:

- Our model is predicting 85% accurately.
- The model performs well on **Class 0**, achieving high recall (0.86), meaning it correctly identifies most of the **Class 0** samples.
- For **Class 1**, performance is poor, with recall only 0.11. This shows the model is failing to detect most of **Class 1** cases.

The model is biased toward **Class 0** and struggles with detecting **Class 1**